

How to: Get Up and Running with Magic WAN (+ its interops)

Erfi Anugrah

2024-05-24

Table of contents

Step 1: Conduit Configuration	2
Step 2: GRE and/or IPsec Tunnel Prerequisites	2
Step 3: Configure your tunnels and static routes on Cloudflare	2
Step 4: Make sure health-checks work	3
Step 5: Route the private subnets to your vti	4
Step 6: Automate your provisioning through the use of Gitops/IaC	5
GRE Tunnels	5
IPsec Tunnels	5
Static Routes	6
Magic Firewall	7
Step 7: Interoperability (Optional)	8
PAC File (Proxy Endpoints)	9
PAC File Worker Deployment	9
Gitops/IaC	13
WARP/Clientless RBI	13
Gitops/IaC	13
Cloudflare Tunnel	16
Gitops/IaC	17
(Beta) WARP Connector	19
Site-to-site Connectivity	19

Step 1: Conduit Configuration

This would be setup during the onboarding with Cloudflare, the setup would require specific information from your end w.r.t specific subnets that should be upgraded or if you want to use non RFC 1918 prefixes.

Step 2: GRE and/or IPsec Tunnel Prerequisites

Both GRE and IPsec would add on top of the raw TCP packet, assuming MTU size being 1500 (this can be lower depending on your internet breakouts, KPN and Deutsch Telecom for instance would already need TCP clamping), the MSS would be lower:

For GRE:

Standard Internet Routable MTU	1500 bytes
- Original IP header	20 bytes
- Original protocol header (TCP)	20 bytes
- New IP header	20 bytes
- New protocol header (GRE)	4 bytes
= Maximum segment size (MSS)	1436 bytes

For IPsec this value would be lower still, depending on if it's IPsec within GRE or on its own. The ESP header would be 8 bytes, ESP Trailer could be 16 to 20 bytes conservatively.

In my case the end value is 1350.

You can check by running this command:

```
1 sudo tcpdump -i [INTERFACE] 'tcp[tcpflags] & (tcp-syn)≠0 and dst port [443 or 80]
```

Or running ping with do-not-fragment and size:

```
1 sudo ping IP -M -s 1500
```

An example would look like this for the internet breakout:

```
1 09:47:29.132389 IP [MY_IP].fixed.kpn.net.53958 > 162.159.138.105.https: Flags [S], seq 471050517, win 64240,
  ↪ options [mss 1452,nop,wscale 8,nop,nop,sackOK], length 0
```

With IPsec, 10.68.100.20 is the tunnel endpoint on my side of the tunnel:

```
1 09:49:19.268016 IP 10.68.100.20.51678 > 104.16.236.133.https: Flags [S], seq 2995991175, win 32120, options
  ↪ [mss 1350,sackOK,TS val 1792385600 ecr 0,nop,wscale 7], length 0:
```

Step 3: Configure your tunnels and static routes on Cloudflare

- Create a [vti](#) with a /31 subnet for use, refer to your vendor documentation on how to create one
- The Cloudflare endpoint will be provided to you via the conduit configuration yaml
- The Customer endpoint would be the IP provided to you via your ISP
- By default, you can only add [static routes](#) with [RFC 1918](#) IP prefixes like:
 - 10.0.0.0/8
 - 172.16.0.0/12
 - 192.168.0.0/16

There are exceptions for publicly routable addresses, inform your friendly (at this point) implementation manager before the project begins.

Step 4: Make sure [health-checks](#) work

You should start seeing Cloudflare's side of the tunnel hitting yours (health wise, take an average as not all data centers pinging your end of the tunnel matters):

```
1 10:05:07.160315 IP 10.68.100.21 > 10.68.100.20: ICMP echo request, id 7295, seq 0, length 64
2 10:05:07.160378 IP 10.68.100.20 > 10.68.100.21: ICMP echo reply, id 7295, seq 0, length 64
3 10:05:07.169088 IP 10.68.100.21 > 10.68.100.20: ICMP echo request, id 63043, seq 0, length 64
4 10:05:07.169150 IP 10.68.100.20 > 10.68.100.21: ICMP echo reply, id 63043, seq 0, length 64
5 10:05:07.208415 IP 10.68.100.21 > 10.68.100.20: ICMP echo request, id 41057, seq 0, length 64
6 10:05:07.208498 IP 10.68.100.20 > 10.68.100.21: ICMP echo reply, id 41057, seq 0, length 64
7 10:05:07.238198 IP 10.68.100.21 > 10.68.100.20: ICMP echo request, id 17771, seq 0, length 64
```

If you have replay-window enabled, make sure [anti-replay protection](#) is enabled on Cloudflare's side.

Step 5: Route the private subnets to your vti

! Important

Check your sysctl configuration for linux kernels.
It is possible that you might need to set these values:

```
1 sudo sysctl -w net.ipv4.conf.all.accept_local=1
2 sudo sysctl -w net.ipv4.conf.all.rp_filter=0
```

For rp_filter:

- 0 - No source validation.
- 1 - Strict mode as defined in RFC3704 Strict Reverse Path Each incoming packet is tested against the FIB and if the interface is not the best reverse path the packet check will fail. By default failed packets are discarded.
- 2 - Loose mode as defined in RFC3704 Loose Reverse Path Each incoming packet's source address is also tested against the FIB and if the source address is not reachable via any interface the packet check will fail.

Current recommended practice in RFC3704 is to enable strict mode to prevent IP spoofing from DDos attacks.

If using asymmetric routing or other complicated routing, then loose mode is recommended.

Consult your vendor documentation, here are some examples:

- [Alibaba Cloud VPN Gateway](#)
- [Amazon AWS Transit Gateway](#)
- [Aruba EdgeConnect Enterprise](#)
- [Cisco IOS XE](#)
- [Cisco SD-WAN](#)
- [Fortinet](#)
- [Furukawa Electric FITELnet](#)
- [Google Cloud VPN](#)
- [Microsoft Azure](#)
- [Palo Alto Networks NGFW](#)
- [pfSense](#)
- [SonicWall](#)
- [Sophos Firewall](#)
- [strongSwan](#)
- [VyOS](#)

You can do a TCP traceroute to see if it's going via the tunnel to another endpoint that's also exposed via Magic WAN:

```
1 sudo traceroute -T 172.18.0.8
2 traceroute to 172.18.0.8 (172.18.0.8), 30 hops max, 60 byte packets
3  1  10.68.69.1 (10.68.69.1)  0.627 ms  0.460 ms  0.480 ms
4  2  10.68.100.21 (10.68.100.21)  5.833 ms  5.289 ms  5.264 ms
5  3  172.71.93.32 (172.71.93.32)  7.644 ms  7.038 ms  10.645 ms
6  4  172.18.0.8 (172.18.0.8)  10.356 ms  10.562 ms  10.244 ms
```

Step 6: Automate your provisioning through the use of Gitops/IaC

This can be done via the UI, API or via Terraform. However, it is recommended to use infrastructure as code as much as possible. The examples make use of `.tfvars` file and a `variables.tf` file to reference those sensitive values when applying or planning the infrastructure with Terraform. There are many ways to go about securing the sensitive information including the `.tfstate` file such as encrypting with [Mozilla SOPS](#) and [Age](#).

GRE Tunnels

```
1 resource "cloudflare_gre_tunnel" "vyos_sg" {
2   account_id      = var.cloudflare_account_id
3   name            = "vyos_sg"
4   customer_gre_endpoint = var.sg_ip
5   cloudflare_gre_endpoint = var.wan_ip_1 # This will be provided to you during onboarding
6   interface_address = "10.68.88.21/31"
7   description      = "vyos_sg_gre"
8   ttl              = 64
9   mtu              = 1476
10  health_check_enabled = true
11  health_check_target  = var.sg_ip
12  health_check_type    = "request"
13 }
```

IPsec Tunnels

```
1 resource "cloudflare_ipsec_tunnel" "vyos_sg_ipsec" {
2   account_id      = var.cloudflare_account_id
3   name            = "vyos_sg_ipsec"
4   customer_endpoint = var.sg_ip
5   cloudflare_endpoint = var.wan_ip_2
6   interface_address = "10.68.77.21/31"
7   description      = "vyos_sg_ipsec_m1"
8   health_check_enabled = true
9   health_check_target  = var.sg_ip
10  health_check_type    = "request"
11  psk                  = var.psk_sg
12  allow_null_cipher    = false
13  hex_id               = var.hex_id_sg
14  fqdn_id              = var.fqdn_id_sg
15  user_id              = var.user_id_sg
16 }
```

Static Routes

```
1 The `next-hop` address would be Cloudflare's side of the tunnel
2 resource "cloudflare_static_route" "eth1_vyos_nl_ipsec" {
3   account_id = var.cloudflare_account_id
4   description = "ETH1"
5   prefix      = "10.68.69.0/24"
6   nexthop     = "10.68.100.20"
7   priority    = 100
8 }
9 resource "cloudflare_static_route" "eth1_100_vyos_nl_ipsec" {
10  account_id = var.cloudflare_account_id
11  description = "VLAN_100"
12  prefix      = "10.68.70.0/24"
13  nexthop     = "10.68.100.20"
14  priority    = 100
15 }
16
17 resource "cloudflare_static_route" "podman_vyos_nl_ipsec" {
18  account_id = var.cloudflare_account_id
19  description = "Podman"
20  prefix      = "172.18.0.0/16"
21  nexthop     = "10.68.100.20"
22  priority    = 100
23 }
```

Magic Firewall

Refer also to [example rulesets](#) based on common attack vectors

```
1 resource "cloudflare_magic_firewall_ruleset" "magic_firewall" {
2   account_id = var.cloudflare_account_id
3   name       = "Magic WAN Firewall"
4   description = "Default Magic WAN Firewall"
5
6   rules = [
7     {
8       action      = "allow"
9       expression  = "(ip.proto eq \"icmp\")"
10      description = "Allow ICMP"
11      enabled     = "true"
12    },
13    {
14      action      = "allow"
15      expression  = "(ip.src in {100.64.0.0/10})"
16      description = "Allow WARP Virtual IPs"
17      enabled     = "true"
18    }
19  ]
20 }
```

Step 7: Interoperability (Optional)

Warning

Be aware of these limitations

Virtual Networks

- Ensure Cloudflare Tunnel and WARP are on the default network for the interoperability to function

DH Group for IPsec Tunnels

- Cloudflare supports DH groups 20, 14, 5, but only one should be used when creating tunnels

WARP w/ Magic WAN

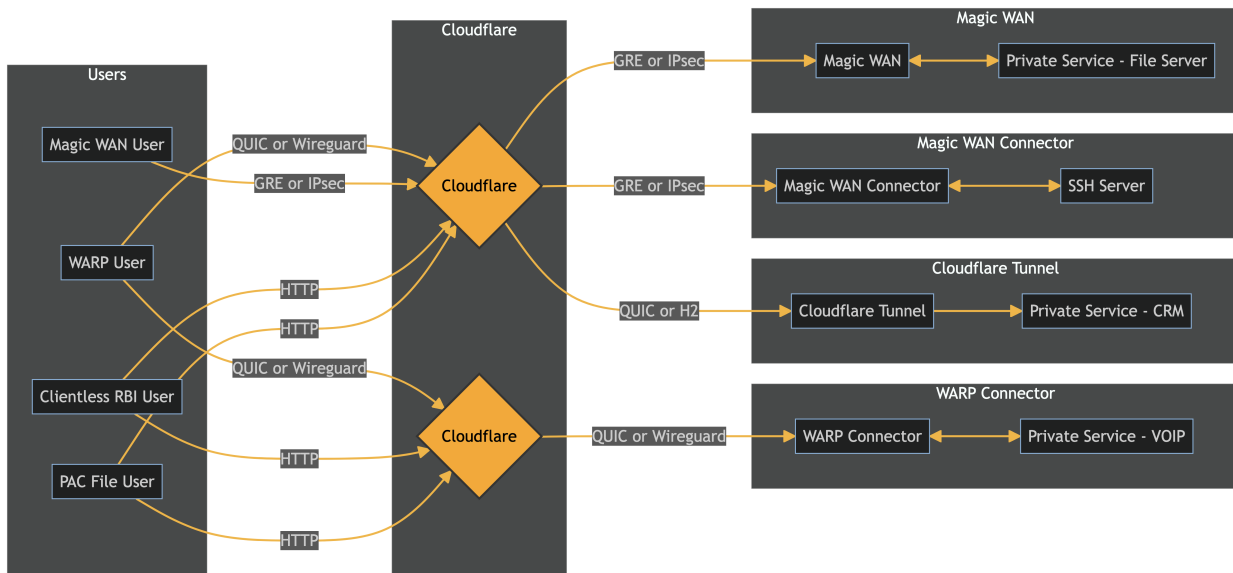
- Ensure WARP connectivity in locations whereby connectivity is routed to Cloudflare Gateway via Magic WAN to be excluded. This is due to the double encapsulation, once by WARP, and again via Magic WAN. Routing policies can be used to exclude connections to the [WARP Ingress IPs](#) and [WARP UDP Ports](#)

Cloudflare Tunnel w/ Magic WAN

- Cloudflare Tunnel does not support outbound connections. Overlapping routes between Cloudflare Tunnel and Magic WAN will cause issues with outbound connections since Cloudflare Tunnel routes are prioritized over Magic WAN static routes

WARP Connector w/ Magic WAN

- The WARP Connector solves the bi-directional use case that Cloudflare Tunnel doesn't solve, however at this point, does not work with Magic WAN when configured within the same Cloudflare account/organisation (as the conduit configuration is tied to the account), and should be used as an alternative (as there are overlapping use cases) if the traditional approach via Magic WAN does not suit your current infrastructure.



1

¹The arrows signify bidirectional connections or unidirectional

PAC File (Proxy Endpoints)

There are many ways to deploy the PAC file such as MDMs and using a remote server, in this case, I'll be using [Workers](#). You can follow the steps [here](#), when using the code, you can create more cases to match specific files, in my case it is nl.pac and sg.pac to simulate the two locations, I also used [Workers Secrets](#) to store the Proxy Endpoint domains:

PAC File Worker Deployment

Worker endpoint (index.js)

```
1 import { nl_pac_file, sg_pac_file } from "./pac_file.js";
2
3 export default {
4   fetch(request, env) {
5     const url = new URL(request.url);
6
7     if (url.pathname === "/nl.pac") {
8       return nl_pac_file(env);
9     } else if (url.pathname === "/sg.pac") {
10      return sg_pac_file(env);
11    } else {
12      return new Response("Not Found", { status: 404 });
13    }
14  },
15 };
```

PAC File variables (pac_file.js)

```
1 export function nl_pac_file(env) {
2   const nl = `
3   function FindProxyForURL(url, host) {
4     // No proxy for private (RFC 1918) IP addresses (intranet sites)
5     // if (
6     //   isInNet(dnsResolve(host), "10.0.0.0", "255.0.0.0") ||
7     //   isInNet(dnsResolve(host), "172.16.0.0", "255.240.0.0") ||
8     //   isInNet(dnsResolve(host), "192.168.0.0", "255.255.0.0")
9     // ) {
10    //   return "DIRECT";
11    // }
12
13    // No proxy for localhost
14    // if (isInNet(dnsResolve(host), "127.0.0.0", "255.0.0.0")) {
15    //   return "DIRECT";
16    // }
17    // Example logic to determine whether to use a proxy
18    return "HTTPS ${env.NL_DOMAIN}.proxy.cloudflare-gateway.com:443";
19  }
20  `;
21  // Set headers to prevent caching
22  const headers = new Headers({
23    "Content-Type": "application/x-ns-proxy-auto-config",
24    "Cache-Control": "no-store, max-age=0",
25  });
26
27  return new Response(nl, { headers: headers });
28 }
29
30 export function sg_pac_file(env) {
31   const sg = `
32   function FindProxyForURL(url, host) {
33     // No proxy for private (RFC 1918) IP addresses (intranet sites)
34     // if (
35     //   isInNet(dnsResolve(host), "10.0.0.0", "255.0.0.0") ||
36     //   isInNet(dnsResolve(host), "172.16.0.0", "255.240.0.0") ||
37     //   isInNet(dnsResolve(host), "192.168.0.0", "255.255.0.0")
38     // ) {
39     //   return "DIRECT";
40     // }
41
42    // No proxy for localhost
43    // if (isInNet(dnsResolve(host), "127.0.0.0", "255.0.0.0")) {
44    //   return "DIRECT";
45    // }
46    // Example logic to determine whether to use a proxy
47    return "HTTPS ${env.SG_DOMAIN}.proxy.cloudflare-gateway.com:443";
48  }
49  `;
```

```
50 // Set headers to prevent caching
51 const headers = new Headers({
52   "Content-Type": "application/x-ns-proxy-auto-config",
53   "Cache-Control": "no-store, max-age=0",
54 });
55
56 return new Response(sg, { headers: headers });
57 }
```

wrangler.toml

```
1 name = "worker-proxy-pac"
2 main = "src/index.js"
3 compatibility_date = "2024-05-12"
4 compatibility_flags = ["nodejs_compat"]
5
6 [dev]
7 port = 9001
8 local_protocol="http"
9 upstream_protocol="https"
10
11 [env.staging]
12 name = "staging-pac"
13 vars = { ENVIRONMENT = "staging" }
14 workers_dev = true
15
16 [env.prod]
17 name = "prod-pac"
18 vars = { ENVIRONMENT = "production" }
19 routes = [
20   { pattern = "proxy.{DOMAIN}.com", custom_domain = true },
21 ]
22 # Secrets
23 # NL_DOMAIN
24 # SG_DOMAIN
```

Gitops/IaC

Proxy Endpoints HCL

```
1 resource "cloudflare_teams_proxy_endpoint" "nl_proxy_endpoint" {
2     account_id = var.cloudflare_account_id
3     name       = "nl"
4     ips        = ["${var.nl_ip}/32"]
5 }
6
7 resource "cloudflare_teams_proxy_endpoint" "sg_proxy_endpoint" {
8     account_id = var.cloudflare_account_id
9     name       = "sg"
10    ips        = ["${var.sg_ip}/32"]
11 }
```

WARP/Clientless RBI

With WARP, the [managed](#) deployment approach is recommended as you can ensure that not only you're managing the devices and the policies but also the WARP client itself. This can also be managed through the use of [Device Profiles](#). If the use case is to connect to endpoints behind Cloudflare Tunnel or Magic WAN, select the default [Virtual Network](#) in the client's dropdown menu.

Both WARP (you can also isolate applications) and [Clientless RBI](#) would use the same Access application and policy as shown below, and both would allow the user to access private IP applications via the other on-ramps.

Gitops/IaC

WARP/RBI Access Policy HCL

```
1 resource "cloudflare_access_policy" "warp_login" {
2     account_id      = var.cloudflare_account_id
3     name            = "Allow Erfi"
4     decision        = "allow"
5     session_duration = "30m"
6
7     include {
8         group = [cloudflare_access_group.erfi_corp.id]
9     }
10 }
```

WARP/RBI Access App HCL

```
1 resource "cloudflare_access_application" "warp_login" {
2   account_id = var.cloudflare_account_id
3   policies = [
4     cloudflare_access_policy.warp_login.id
5   ]
6   allowed_idps = [
7     cloudflare_access_identity_provider.entra_id.id,
8     cloudflare_access_identity_provider.google_workspace.id,
9     cloudflare_access_identity_provider.gmail.id,
10    cloudflare_access_identity_provider.keycloak_oidc.id,
11    cloudflare_access_identity_provider.authentic_oidc.id,
12    cloudflare_access_identity_provider.authentic_saml.id,
13    cloudflare_access_identity_provider.otp.id
14  ]
15  auto_redirect_to_identity = false
16  domain                    = "erfianugrah.cloudflareaccess.com/warp"
17  name                     = "Warp Login App"
18  session_duration         = "24h"
19  type                     = "warp"
20 }
```

Virtual Networks HCL

```
1 resource "cloudflare_tunnel_virtual_network" "vyos_nl" {
2   account_id = var.cloudflare_account_id
3   name       = "vyos_nl_vnet"
4   is_default_network = true
5 }
```

Device Settings Policy HCL

```
1 resource "cloudflare_device_settings_policy" "default" {
2   account_id      = var.cloudflare_account_id
3   name            = "default"
4   description     = "default_policy"
5   # precedence    = 100
6   # match         = "any(identity.groups.name[*] in {\\"Erfi Corp\\"})"
7   default        = true
8   enabled        = true
9   allow_mode_switch = true
10  allow_updates   = true
11  allowed_to_leave = true
12  auto_connect   = 0
13  captive_portal  = 180
14  disable_auto_fallback = false
15  switch_locked  = false
16  service_mode_v2_mode = "warp"
17  service_mode_v2_port = 3000
18  exclude_office_ips = true
19 }
20
21 resource "cloudflare_device_settings_policy" "google" {
22   account_id      = var.cloudflare_account_id
23   name            = "Google Workspace"
24   description     = "google_workspace_policy"
25   precedence      = 200
26   match          = "any(identity.groups.name[*] in {\\"Erfi Corp\\"})"
27   default        = false
28   enabled        = true
29   allow_mode_switch = true
30   allow_updates   = true
31   allowed_to_leave = true
32   auto_connect   = 0
33   captive_portal  = 180
34   disable_auto_fallback = false
35   switch_locked  = false
36   service_mode_v2_mode = "warp"
37   service_mode_v2_port = 3000
38   exclude_office_ips = true
39 }
```

Cloudflare Tunnel

Tunnels are used to exposed private applications or for connectivity in the case of RDP, SSH, VNC and the like, it does not support bi-directional traffic as mentioned above. By setting up the routes, you can now reach those same private applications behind Cloudflare Tunnel be it from a PAC file, clientless RBI, WARP or behind Magic WAN.

Note

Cloudflare Tunnel Deployment

There are (again) like WARP many ways to deploy, as systemd on a VM or metal, docker standalone or in compose or a deployment in k8s or docker swarm.

Be aware of the system requirements:

For most use cases, we recommend the following baseline configuration:

- Run a cloudflared replica on two dedicated host machines per network location. Using two hosts enables server-side redundancy and traffic balancing.
- Size each host with minimum 4GB of RAM and 4 CPU cores.
- Allocate 50,000 ports to the cloudflared process on each host.

This setup is usually sufficient to handle traffic from 8,000 WARP users (4,000 per host). The actual amount of resources used by cloudflared will depend on many variables, including the number of requests per second, bandwidth, network path and hardware. As additional users are onboarded, or if network traffic increases beyond your existing tunnel capacity, you can scale your tunnel by adding an additional cloudflared host in that location.

Gitops/IaC

Terraform provider HCL with random provider

```
1 terraform {
2   required_providers {
3     cloudflare = {
4       source = "cloudflare/cloudflare"
5       version = "~> 4.0"
6     }
7     random = {
8       source = "hashicorp/random"
9       version = "~> 3.0"
10    }
11  }
12 }
```

Use the random provider to generate string that will be use to set the tunnel secret

```
1 # Generate a random string
2 resource "random_string" "tunnel_secret" {
3   length = 32
4   special = false
5 }
```

Use the random string as the tunnel secret to create the Cloudflare Tunnel

```
1 resource "cloudflare_tunnel" "vyos_n1" {
2   account_id = var.cloudflare_account_id
3   name       = "vyos_n1"
4   secret     = base64encode(random_string.tunnel_secret.result)
5   config_src = "cloudflare"
6 }
```

Warning

Again, be aware of conflicting routes between Cloudflare Tunnel and Magic WAN

Tunnel Route HCL

```
1 resource "cloudflare_tunnel_route" "vyos_nl" {
2   account_id      = var.cloudflare_account_id
3   tunnel_id       = cloudflare_tunnel.vyos_nl.id
4   network         = "0.0.0.0/0"
5   virtual_network_id = cloudflare_tunnel_virtual_network.vyos_nl.id
6 }
```

Tunnel Config HCL

```
1 resource "cloudflare_tunnel_config" "vyos_nl" {
2   account_id = var.cloudflare_account_id
3   tunnel_id  = cloudflare_tunnel.vyos_nl.id
4
5   config {
6     warp_routing {
7       enabled = true
8     }
9     ingress_rule {
10      hostname = "prom-tunnel-nl.${var.domain_name}"
11      service  = "http://localhost:11000"
12    }
13    ingress_rule {
14      hostname = "prom-caddy-nl.${var.domain_name}"
15      service  = "http://172.18.0.4:2018"
16    }
17    ingress_rule {
18      service = "http_status:404"
19    }
20  }
21 }
```

(Beta) WARP Connector

It shares the same functionality as Cloudflare Tunnel does but with bidirectional support. This means that only you can expose private services/applications which require two-way traffic such as communication services.

As the WARP connector is also essentially a device connected via WARP to the internet, a remote user with WARP connectivity can also connect to it directly and not just the services exposed behind it which is what we call WARP-to-WARP connectivity. Both site-to-site and peer-to-peer use cases work here, similar to Magic WAN.

Site-to-site Connectivity

1. Setup Access policy for WARP enrolment (this would be the same policy for WARP authentication), since it's site-to-site, we will have to deploy it via `.xml` with the required parameters.

Refer to [service tokens](#) and [device enrolment](#) for details.

2. Setup the split tunnel to either include or exclude mode, in this example, I have chosen exclude so that I can still SSH into the nodes to showcase the functionality.

Since this portion can be managed by Terraform, this is the example HCL:

```
1 resource "cloudflare_split_tunnel" "default_include" {
2   account_id = var.cloudflare_account_id
3   policy_id   = cloudflare_device_settings_policy.default.id
4   mode        = "exclude"
5   tunnels {
6     address    = "10.68.69.3"
7     description = "ERFI1"
8   }
9   tunnels {
10    address    = "10.68.73.3"
11    description = "arch-0"
12  }
13  tunnels {
14    address    = "10.68.73.2"
15    description = "pve"
16  }
17 }
```

3. Ensure these settings are enabled:

- Enable CGNAT routing (Settings -> Network)
 - Enable Proxy (UDP, TCP, ICMP)
 - Enable WARP to WARP
 - Override local interface IP (Settings -> WARP Client)
 - * If in include mode, ensure 100.96.0.0/12 is included

Some of these settings are available as a Terraform resource (as shown below), since it would be a mix a of API/UI and Terraform, Terraform would not know the state, so just keep to API/UI for all settings.

```
1 resource "cloudflare_teams_account" "miau3" {
2   account_id           = var.cloudflare_account_id
3   tls_decrypt_enabled  = false
4   protocol_detection_enabled = false
5   activity_log_enabled = true
6   non_identity_browser_isolation_enabled = true
7
8   body_scanning {
9     inspection_mode = "deep"
10  }
11
12  antivirus {
13    enabled_download_phase = false
14    enabled_upload_phase   = false
15    fail_closed            = false
16  }
17
18  fips {
19    tls = false
20  }
21
22  proxy {
23    tcp   = true
24    udp   = true
25    root_ca = true
26  }
27
28  url_browser_isolation_enabled = true
29
30  logging {
31    redact_pii = false
32    settings_by_rule_type {
33      dns {
34        log_all    = true
35        log_blocks = false
36      }
37      http {
38        log_all    = true
39        log_blocks = false
40      }
41    }
42  }
```

```

42     log_all    = true
43     log_blocks = false
44   }
45 }
46 }
47
48 extended_email_matching {
49   enabled = true
50 }
51 }

```

4. Set up the WARP Connector

- Either via API or the UI, create the WARP connector (currently not available as a Terraform resource):

```

1 curl --request POST \
2   --url https://api.cloudflare.com/client/v4/accounts/account_id/warp_connector \
3   --header 'Content-Type: application/json' \
4   --header 'X-Auth-Email: ' \
5   --data '{
6     "name": "arch-0"
7   }'

```

Upon creation, you will get the example .xml e.g. which will look like this:

```

1 <dict>
2   <key>organization</key>
3   <string>org_name</string>
4   <key>auth_client_id</key>
5   <string>client_access_id_string</string>
6   <key>auth_client_secret</key>
7   <string>client_access_secret_string</string>
8   <key>warp_connector_token</key>
9   <string>connector_token_string</string>
10 </dict>

```

This is what you will place in the /var/lib/cloudflare-warp directory after the WARP client install. Refer to the relevant [repositories](#) that matches your Linux distribution, for Arch, you would have to install via [AUR](#).

Once done, if not yet enabled, run `sudo systemctl enable warp-svc` then `sudo systemctl start warp-svc`, if already running as a service, run `sudo systemctl restart warp-svc` so that it will use the new configuration to connect to Cloudflare.

5. Set up Linux networking

We need to enable IP forwarding and MSS clamping (WARP's MTU is 1280 bytes).

Create a .conf file in the etc/sysctl.d/ directory with value of net.ipv4.ip_forward = 1 so that this persists upon reboot.

To check what IPtables rules are existing, we can run `sudo iptables-save > test.conf` to just see what's currently being use, if all's well. We can add the following to the existing /etc/iptables/iptables.rules file, such that it would look something like this:

```
1 *filter
2 :INPUT ACCEPT [0:0]
3 :FORWARD ACCEPT [0:0]
4 :OUTPUT ACCEPT [0:0]
5 -A FORWARD -i CloudflareWARP -p tcp -m tcp --tcp-flags SYN,RST SYN -j TCPMSS --clamp-mss-to-pmtu
6 -A FORWARD -o CloudflareWARP -p tcp -m tcp --tcp-flags SYN,RST SYN -j TCPMSS --clamp-mss-to-pmtu
7 COMMIT
```

This would make sure the MSS clamping rules will also persist on reboot.

6. Route nodes/VMs/containers to the WARP connector

The most common use case here would either be an alternate gateway to the main router or an intermediate gateway. In this example, I will be using it as an intermediate gateway.

Set up the static routes for each WARP connector, the configuration will be similar to the Cloudflare Tunnel routes. In our case, on the WARP connector called arch-1, we will exposing 10.68.73.102/32 for the arch-2 machine, and arch-3 will be exposing 10.68.73.104/32 for arch-4 machine.

Let's use arch-2 as an example as the same steps will be replicated on the other node.

Run:

```
1 ip route
```

and you should be able to see default routes, for instance:

```
1 default via 10.68.73.1 dev ens18 proto dhcp src 10.68.73.102 metric 100
```

we want to create another route that has higher precedence than this (lower number).

Run:

```
1 sudo ip route add default via 10.68.73.101 dev ens18 metric 99
```

What this means is that, traffic from arch-2 will go via the ens18 interface to 10.68.73.101 which is the WARP connector arch-1 and using it as a gateway, now this node is connected to Cloudflare.

Do the same for the other node, and change the IP.

Run:

```
1 curl https://icanhazip.com
```

on either machine, and you should be able to see a Cloudflare IP, just be aware since all traffic is sent to the WARP connector, that the DNS server used by the machines routed to WARP connector can be accessed while on the network, either by another route or a public DNS server.

```
1 curl icanhazip.com
2 104.28.218.39
```

7. Route the machines to each other

We can either router the entire prefix to the WARP connector or specific IPs, in my case, I will be using just one IP:

```
1 sudo ip route add 10.68.73.104 via 10.68.73.101 dev ens18 metric 98
```

This will allow 10.68.73.102 to reach 10.68.73.104, do the same on arch-4, just by swapping the IPs, from .104 to .102

Run:

```
1 ping 10.68.73.102
```

or

```
1 mtr 10.68.73.102
```

and vice versa to see the traffic reach the WARP connector CGNAT IP before reaching the other machine.